

Problem 1

Calculate the Lagrangian $\mathcal{L}(x, \alpha)$.

$$\mathcal{L}(x, \alpha) = -(x_1 + x_2) + \alpha_1(x_1^2 + x_2^2 - 1)$$

Obtain the Lagrangian dual function $g(\alpha)$.

$$\begin{aligned}\frac{\nabla \mathcal{L}(x, \alpha)}{\partial x_1} &= 0 \Leftrightarrow \\ 2\alpha_1 x_1 - 1 &= 0 \Leftrightarrow \\ x_1 &= \frac{1}{2\alpha_1}\end{aligned}$$

$$\begin{aligned}\frac{\nabla \mathcal{L}(x, \alpha)}{\partial x_2} &= 0 \Leftrightarrow \\ 2\alpha_1 x_2 - 1 &= 0 \Leftrightarrow \\ x_2 &= \frac{1}{2\alpha_1}\end{aligned}$$

Solve the dual problem (Plug x^* in $g(\alpha)$).

$$\begin{aligned}\frac{\nabla g(\alpha)}{\partial \alpha} &= 0 \Leftrightarrow \\ \frac{1}{\alpha_1^2} - \frac{1}{2\alpha_1^2} - 1 &= 0 \Leftrightarrow \\ \alpha_1^2 &= \frac{1}{2} && \text{(by constraint } \alpha_i \geq 0) \\ \alpha_1 &= \frac{1}{\sqrt{2}}\end{aligned}$$

Problem 2

- Similarities
 - Both algorithms try to solve the problem of binary classification by finding a decision boundary $w^\top x + b = 0$ that separates all datapoints x_i with label 1 from all datapoints x_j with label -1
 - Differences
 - SVM has a closed form solution
 - SVM gives an unique solution (constrained optimization) by choosing a decision boundary s.t. it has a maximum margin to its nearest datapoints
 - Perceptron must be solved iteratively
 - Perceptron may have infinitely many correct solutions (if available) (unconstrained optimization)
-

Problem 3

By the formulation of the SVM problem we have

$$\begin{aligned} & \text{minimize} \quad f_0(w, b) = \frac{1}{2}w^\top w \\ & \text{subject to} \quad f_i(w, b) = y_i(w^\top x_i + b) - 1 \geq 0, \quad \text{for } i = 1, \dots, N \end{aligned} \quad (1)$$

Clearly we can rewrite (1) to

$$f_i(w, b) = -y_i(w^\top x_i + b) + 1 \leq 0, \quad \text{for } i = 1, \dots, N$$

By Slater's constraint qualification we have that the duality gap of the SVM problem is zero if $f_0(x)$, $f_1(x)$, ... $f_N(x)$ are convex and the constraints $f_1(x)$, ... $f_N(x)$ are affine. Clearly both assumptions are met, because $f_0(x)$ is simply the L_2 -norm, which is convex and the constraints are linear functions in w shifted by an offset b , which makes them affine. Thus the duality gap is zero.

Problem 4

a). Let $\mathbf{X} \in \mathbb{R}^{n \times d}$, where each row is a datapoint $\mathbf{x} \in \mathbb{R}^d$, $\mathbf{Y} \in \mathbb{R}^{n \times n}$, where each column is the label y_i of the i -th datapoint replicated n times and $\boldsymbol{\alpha} \in \mathbb{R}^n$ with α_i at position i .

By the Hadamard product we have

$$\mathbf{Q} = -\mathbf{X}\mathbf{X}^\top \odot (\mathbf{Y} \odot \mathbf{Y}^\top) \quad (1)$$

$$g(\boldsymbol{\alpha}) = \boldsymbol{\alpha}^\top \mathbf{Q} \boldsymbol{\alpha} \quad (2)$$

b). We can reformulate (2) to $Q = -(\mathbf{X} \odot \mathbf{Y})^\top (\mathbf{X} \odot \mathbf{Y})$ and define $A = (\mathbf{X} \odot \mathbf{Y})$. By construction we have $A^\top A$ is positive semi-definite, i.e. $\forall z : z^\top A^\top A z \geq 0$ because $z^\top A^\top A z = (Az)^\top A z \geq 0$ (i.e. L_2 norm is non-negative). Thus $Q = -A^\top A$ is negative semi-definite.

c). From negative semi-definiteness of Q and the Hessian of $\boldsymbol{\alpha}^\top Q \boldsymbol{\alpha}$ is negative, it follows that $\boldsymbol{\alpha}^\top Q \boldsymbol{\alpha}$ is a concave function. Since in the dual formulation we maximize $g(\boldsymbol{\alpha})$, $\frac{\nabla g(\boldsymbol{\alpha})}{\partial \boldsymbol{\alpha}} = 0$ is a sufficient condition to get the global maximum $\boldsymbol{\alpha}^*$.

07_homework_svm-2

December 10, 2017

1 Programming assignment 7: SVM

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

from sklearn.datasets import make_blobs

from cvxopt import matrix, solvers
```

1.1 Your task

In this sheet we will implement a simple binary SVM classifier.

We will use CVXOPT <http://cvxopt.org/> - a Python library for convex optimization. If you use Anaconda, you can install it using


```
conda install cvxopt
```

As usual, your task is to fill out the missing code, run the notebook, convert it to PDF and attach it to your HW solution.

1.2 Generate and visualize the data

```
In [2]: N = 100 # number of samples
D = 2 # number of dimensions
C = 2 # number of classes
seed = 3 # for reproducible experiments

X, y = make_blobs(n_samples=N, n_features=D, centers=2, random_state=seed)
y[y == 0] = -1 # it is more convenient to have {-1, 1} as class labels (instead of {0, 1})
y = y.astype(np.float)
plt.figure(figsize=[10, 8])
plt.scatter(X[:, 0], X[:, 1], c=y)
plt.show()
```



1.3 Task 1: Solving the SVM dual problem

Remember, that the SVM dual problem can be formulated as a Quadratic programming (QP) problem. We will solve it using a QP solver from the CVXOPT library.

The general form of a QP is

$$\min_{\mathbf{x}} \frac{1}{2} \mathbf{x}^T \mathbf{P} \mathbf{x} + \mathbf{q}^T \mathbf{x}$$

$$\text{subject to } \mathbf{G} \mathbf{x} \preceq \mathbf{h}$$

$$\text{and } \mathbf{A} \mathbf{x} = \mathbf{b}$$

where \preceq denotes "elementwise less than or equal to".

Your task is to formulate the SVM dual problems as a QP and solve it using CVXOPT, i.e. specify the matrices \mathbf{P} , \mathbf{G} , \mathbf{A} and vectors \mathbf{q} , \mathbf{h} , \mathbf{b} .

```
In [3]: def solve_dual_svm(X, y):
        """Solve the dual formulation of the SVM problem.

        Parameters
```

```

-----
X : array, shape [N, D]
    Input features.
y : array, shape [N]
    Binary class labels (in {-1, 1} format).

Returns
-----
alphas : array, shape [N]
    Solution of the dual problem.
"""
# TODO
# These variables have to be of type cvxopt.matrix
P = matrix(np.dot(X, X.T) * y[None, :] * y[:, None])
q = matrix(-np.ones(shape=(N, 1)))
G = matrix(-np.eye(N))
h = matrix(-np.zeros(shape=(N,1)))
A = matrix(y[None, :])
b = matrix(0.0)

solvers.options['show_progress'] = False
solution = solvers.qp(P, q, G, h, A, b)
alphas = np.array(solution['x'])
return alphas

```

1.4 Task 2: Recovering the weights and the bias

```

In [4]: def compute_weights_and_bias(alpha, X, y):
        """Recover the weights w and the bias b using the dual solution alpha.

        Parameters
        -----
        alpha : array, shape [N]
            Solution of the dual problem.
        X : array, shape [N, D]
            Input features.
        y : array, shape [N]
            Binary class labels (in {-1, 1} format).

        Returns
        -----
        w : array, shape [D]
            Weight vector.
        b : float
            Bias term.
        """
        w = np.dot(X.T, alpha * y[:, None])
        idx_sv = np.where(alpha > 1e-4)[0][0]

```

```

b = y[idx_sv] - np.dot(X[idx_sv, :], w)
return w, b

```

1.5 Visualize the result (nothing to do here)

```

In [5]: def plot_data_with_hyperplane_and_support_vectors(X, y, alpha, w, b):
        """Plot the data as a scatter plot together with the separating hyperplane.

        Parameters
        -----
        X : array, shape [N, D]
            Input features.
        y : array, shape [N]
            Binary class labels (in {-1, 1} format).
        alpha : array, shape [N]
            Solution of the dual problem.
        w : array, shape [D]
            Weight vector.
        b : float
            Bias term.
        """
        plt.figure(figsize=[10, 8])
        # Plot the hyperplane
        slope = -w[0] / w[1]
        intercept = -b / w[1]
        x = np.linspace(X[:, 0].min(), X[:, 0].max())
        plt.plot(x, x * slope + intercept, 'k-', label='decision boundary')
        # Plot all the datapoints
        plt.scatter(X[:, 0], X[:, 1], c=y)
        # Mark the support vectors
        support_vecs = (alpha > 1e-4).reshape(-1)
        plt.scatter(X[support_vecs, 0], X[support_vecs, 1], c=y[support_vecs], s=250, marker='x')
        plt.xlabel('$x_1$')
        plt.ylabel('$x_2$')
        plt.legend(loc='upper left')

```

The reference solution is

```

w = array([[ -0.69192638],
          [-1.00973312]])

```

```

b = 0.907667782

```

Indices of the support vectors are

```

[38, 47, 92]

```

```


In [9]: alpha = solve_dual_svm(X, y)
        w, b = compute_weights_and_bias(alpha, X, y)

```

```

plot_data_with_hyperplane_and_support_vectors(X, y, alpha, w, b)
plt.show()
print('weights w = \n {} \n bias b = \n {}'.format(w, b))
print('Indices of the support vectors are {}'.format(np.where(alpha > 1e-4)[0]))

```



```

weights w =
  [[-0.69192638]
   [-1.00973312]]
bias b =
  [ 0.90766782]
Indices of the support vectors are [38 47 92]

```